

Project 0: Getting Real

陈震雄

武汉大学

2025 年 3 月 15 日



Table of Contents

- 1 Task 1: Booting Pintos
- 2 Task 2: Debugging
- 3 Task 3: Kernel Monitor



Table of Contents

1 Task 1: Booting Pintos

2 Task 2: Debugging

3 Task 3: Kernel Monitor



Exercise 1.1

- Take screenshots of the successful booting of Pintos in QEMU and Bochs.

```
root@6d961ed8cf68:~/pintos/src/threads/build# pintos --
qemu-system-i386 -device isa-debug-exit -drive format=raw,media=disk,index=0,file=/tmp/gj8XLiLta9.dsk -m 4 -net none -nographic -monitor null
Pintos hda1
Loading.....
Kernel command line:
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 104,755,200 loops/s.
Boot complete.

root@6d961ed8cf68:~/pintos/src/threads/build# pintos --bochs --
squish-pty bochs -q
=====
Bochs x86 Emulator 2.6.2
Built from SVN snapshot on May 26, 2013
Compiled on Mar 1 2022 at 16:09:16
=====
00000000000i[ ] reading configuration from bochsrc.txt
00000000000e[ ] bochsrc.txt:8: 'user_shortcut' will be replaced by new 'keyboard' option.
00000000000i[ ] installing nogui module as the Bochs GUI
00000000000i[ ] using log file bochsout.txt
Pintos hda1
Loading.....
Kernel command line:
Pintos booting with 4,096 kB RAM...
383 pages available in kernel pool.
383 pages available in user pool.
Calibrating timer... 102,400 loops/s.
Boot complete.
```



Table of Contents

1 Task 1: Booting Pintos

2 Task 2: Debugging

3 Task 3: Kernel Monitor



Exercise 2.1

- What is the first instruction that gets executed?

```
1 (gdb) debugpintos
2 The target architecture is assumed to be i8086
3 [f000:fff0] 0xffff0: ljmp $0x3630,$0xf000e05b
4 0x0000fff0 in ?? ()
```



Exercise 2.1

+-----+	<- 0xFFFFFFFF (4GB)
32-bit	
memory mapped	
devices	
 /\/\/\/\/\/\/\/\/\	
 /\/\/\/\/\/\/\/\/\	
Unused	
+-----+	<- depends on amount of RAM
 Extended Memory	
+-----+	<- 0x00100000 (1MB)
BIOS ROM	
+-----+	<- 0x000F0000 (960KB)
16-bit devices,	
expansion ROMs	
+-----+	<- 0x000C0000 (768KB)
VGA Display	
+-----+	<- 0x000A0000 (640KB)
 Low Memory	
+-----+	<- 0x00000000

Exercise 2.2

- How does the bootloader read disk sectors? In particular, what BIOS interrupt is used?

```
1 (gdb) break *0x7c00
2 Breakpoint 1 at 0x7c00
3 (gdb) c
4 Continuing.
5 [ 0:7c00] => 0x7c00:  sub    %eax,%eax
6 Breakpoint 1, 0x00007c00 in ?? ()
7 (gdb) x/8i $pc
8 => 0x7c00:      sub    %eax,%eax
9               0x7c02:  mov    %eax,%ds
10              0x7c04:  mov    %eax,%ss
11              0x7c06:  mov    $0xf000,%sp
12              0x7c0a:  add    %al,(%eax)
13              0x7c0c:  sub    %edx,%edx
14              0x7c0e:  mov    $0xe3,%al
15              0x7c10:  int    $0x14
16              ...
17              call read_sector
```


Exercise 2.2

loader.S

```
1 read_sector:
2     pusha
3     sub %ax, %ax
4     push %ax          # LBA sector number [48:63]
5     push %ax          # LBA sector number [32:47]
6     push %ebx         # LBA sector number [0:31]
7     push %es          # Buffer segment
8     push %ax          # Buffer offset (always 0)
9     push $1           # Number of sectors to read
10    push $16           # Packet size
11    mov $0x42, %ah     # Extended read
12    mov %sp, %si       # DS:SI -> packet
13    int $0x13          # Error code in CF
14    popa               # Pop 16 bytes, preserve flags
```



Exercise 2.2

Reading hard disk sectors requires the use of the functions provided by BIOS. Specifically, as mentioned in the title, it triggers a BIOS interrupt. The instruction is located in line 242 (red box in the figure). A complete interrupt table is available for query under the BIOS interrupt call entry on Wikipedia:

13h	0Dh	Reset Fixed Disk Controller
	15h	Get Drive Type
	16h	Get Floppy Drive Media Change Status
	17h	Set Disk Type
	18h	Set Floppy Drive Media Type
	41h	Extended Disk Drive (EDD) Installation Check
	42h	Extended Read Sectors
	43h	Extended Write Sectors
	44h	Extended Verify Sectors
	45h	Lock/Unlock Drive
	46h	Eject Media



Exercise 2.2

- How does the bootloader decide whether it successfully finds the Pintos kernel?



Exercise 2.2

```
1  # Check for MBR signature--if not present, it's not a
2  # partitioned hard disk.
3  cmpw $0xaa55, %es:510
4  jne next_drive
5
6  mov $446, %si    # Offset of partition table entry 1.
7  mov $'1', %al
8  check_partition:
9  # Is it an unused partition?
10  cmpl $0, %es:(%si)
11  je next_partition
12  # Print [1-4].
13  call putc
14
15  # Is it a Pintos kernel partition?
16  cmpb $0x20, %es:4(%si)
17  jne next_partition
18
19  # Is it a bootable partition?
20  cmpb $0x80, %es:(%si)
21  je load_kernel
```

Exercise 2.2

```
1 next_partition:
2   # No match for this partition, go on to the next one.
3   add $16, %si    # Offset to next partition table entry.
4   inc %al
5   cmp $510, %si
6   jb check_partition
7 next_drive:
8   # No match on this drive, go on to the next one.
9   inc %dl
10  jnc read_mbr
```

Element (offset)	Size	Description
0	byte	Bitflags field: 1 = not bootable, 0x81 = bootable (or "active")
1	byte	Signature-1 (0x14)
2	uint16_t	Partition Start LBA (high 16-bit of 48 bit value)
4	byte	System ID
5	byte	Signature-2 (0xeb)
6	uint16_t	Partition Length (high 16-bit of 48 bit value)
8	uint32_t	Partition Start LBA (low uint32_t)
12	uint32_t	Partition Length (low uint32_t)

Exercise 2.2

- What happens when the bootloader could not find the Pintos kernel?

```
1 no_such_drive:
2 no_boot_partition:
3   # Didn't find a Pintos kernel partition anywhere, give up.
4   call puts
5   .string "\rNot found\r"
6
7   # Notify BIOS that boot failed.  See [IntrList].
8   int $0x18
```



Exercise 2.2

18h

Execute [Cassette BASIC](#): On IBM machines up to the early PS/2 line, this interrupt would start the ROM Cassette BASIC. Clones did not have this feature and different machines/BIOSes would perform a variety of different actions if INT 18h was executed, most commonly an error message stating that no bootable disk was present. Modern machines would attempt to [boot from a network](#) through this interrupt. On modern machines this interrupt will be treated by the BIOS as a signal from the bootloader that it failed to complete its task. The BIOS can then take appropriate next steps.^[3]



Exercise 2.2

- At what point and how exactly does the bootloader transfer control to the Pintos kernel?



Exercise 2.2

loader.S

```
1 load_kernel:
2     ...
3     mov $0x2000, %ax
4     mov %ax, %es
5     mov %es:0x18, %dx
6     mov %dx, start
7     movw $0x2000, start + 2
8     ljmp *start
```

0x18	4	8	e_entry	This is the memory address of the entry point from where the process starts executing. This field is either 32 or 64 bits long, depending on the format defined earlier (byte 0x04). If the file doesn't have an associated entry point, then this holds zero.
------	---	---	---------	--



Exercise 2.3

- At the entry of `pintos_init()`, what is the value of the expression `init_page_dir[pd_no(ptov(0))]` in hexadecimal format?



Exercise 2.3

```
1 (gdb) b pintos_init
2 Breakpoint 1 at 0xc00202b6: file ../../threads/init.c, line 78.
3 (gdb) continue
4 Continuing.
5 The target architecture is assumed to be i386
6 => 0xc00202b6 <pintos_init>:      push    %ebp
7
8 Breakpoint 1, pintos_init () at ../../threads/init.c:78
9 (gdb) p init_page_dir[pd_no(ptov(0))]
10 => 0xc000efef:  int3
11 => 0xc000efef:  int3
12 $1 = 0
```



Exercise 2.3

- When `palloc_get_page()` is called for the first time,
 - what does the call stack look like?
 - what is the return value in hexadecimal format?
 - what is the value of expression `init_page_dir[pd_no(ptov(0))]` in hexadecimal format?



Exercise 2.3

```
1 (gdb) b palloc_get_page
2 Breakpoint 2 at 0xc002311a: file ../../threads/palloc.c, line 113.
3 (gdb) continue
4 Continuing.
5 => 0xc002311a <palloc_get_page+6>:      sub    $0x8,%esp
6 Breakpoint 2, palloc_get_page (flags=(PAL_ASSERT | PAL_ZERO))
7 at ../../threads/palloc.c:113
8 (gdb) bt
9 #0  palloc_get_page (flags=(PAL_ASSERT | PAL_ZERO)) at
10 ../../threads/palloc.c:113
11 #1  0xc00203aa in paging_init () at ../../threads/init.c:168
12 #2  0xc002031b in pintos_init () at ../../threads/init.c:100
13 #3  0xc002013d in start () at ../../threads/start.S:180
14 (gdb) fin
15 Run till exit from #0  palloc_get_page (flags=(PAL_ASSERT | PAL_ZERO)) at
16 ../../threads/palloc.c:113
17 => 0xc00203aa <paging_init+17>: add    $0x10,%esp
18 0xc00203aa in paging_init () at ../../threads/init.c:168
19 Value returned is $2 = (void *) 0xc0101000
20 (gdb) p/x init_page_dir[pd_no(ptov(0))]
21 => 0xc000ef8f:  int3
22 => 0xc000ef8f:  int3
23 $3 = 0x0
```



Exercise 2.3

- When `palloc_get_page()` is called for the third time,
 - what does the call stack look like?
 - what is the return value in hexadecimal format?
 - what is the value of expression `init_page_dir[pd_no(ptov(0))]` in hexadecimal format?



Exercise 2.3

```
1 (gdb) bt
2 #0  palloc_get_page (flags=PAL_ZERO) at ../../threads/palloc.c:113
3 #1  0xc0020a81 in thread_create (name=0xc002e895 "idle", priority=0,
4 function=0xc0020eb0 <idle>, aux=0xc000efbc) at ../../threads
5 /thread.c:178
6 #2  0xc0020976 in thread_start () at ../../threads/thread.c:111
7 #3  0xc0020334 in pintos_init () at ../../threads/init.c:119
8 #4  0xc002013d in start () at ../../threads/start.S:180
9 (gdb) fin
10 Run till exit from #0  palloc_get_page (flags=PAL_ZERO) at
11 ../../threads/palloc.c:113
12 => 0xc0020a81 <thread_create+55>:      add    $0x10,%esp
13 0xc0020a81 in thread_create (name=0xc002e895 "idle", priority=0,
14 function=0xc0020eb0 <idle>, aux=0xc000efbc) at ../../
15 /threads/thread.c:178
16 Value returned is $4 = (void *) 0xc0103000
17 (gdb) p/x init_page_dir[pd_no(ptov(0))]
18 => 0xc000ef4f:  int3
19 => 0xc000ef4f:  int3
20 $5 = 0x102027
```

Table of Contents

- 1 Task 1: Booting Pintos
- 2 Task 2: Debugging
- 3 Task 3: Kernel Monitor



Exercise 3.1

- Enhance `threads/init.c` to implement a tiny kernel monitor in Pintos.
Requirements:
 - It starts with a prompt `WHUOS>` and waits for user input.
 - As the user types in a printable character, display the character.
 - When a newline is entered, it parses the input and checks if it is `whoami`. If it is `whoami`, print your student id. Afterward, the monitor will print the command prompt `WHUOS>` again in the next line and repeat.
 - If the user input is `exit`, the monitor will quit to allow the kernel to finish. For the other input, print invalid command. Handling special input such as backspace is not required.
 - If you implement such an enhancement, mention this in your design document.



Exercise 3.1

```
size_t max_len = 10;
char* buf = (char*) malloc(max_len);
while(1)
{
    printf("WHUOS> ");
    memset(buf, '\\0', max_len);
    size_t index = 0;
    while(1)
    {
        char c = input_getc();
        if (c == 13)
        {
            printf("\\n");
            break;
        }
        if (c == 127)
        {
            if (index > 0) {
                buf[--index] = '\\0';
                printf("\\b \\b");
            }
        }
    }
}
```

```
        continue;
    }
    if (index >= max_len) continue;
    buf[index++] = c;
    if (c > 31 && c < 127)
    {
        printf("%c", c);
    }
    if (!strcmp(buf, "whoami"))
    {
        printf("20250227\\n");
        continue;
    }
    if (!strcmp(buf, "exit"))
        break;
    printf("invalid command\\n");
}
free(buf);
printf("Bye!");
}
```

Exercise 3.1

```
root@c9803698e3d0:~/pintos/src/threads# pintos --
qemu-system-i386 -device isa-debug-exit -drive format=raw,media=disk,index=0,file=/tmp/Y8EXWN8C3H.dsk -m 4 -net none -nographic -monitor null
Pintos hda1
Loading.....
Kernel command line:
Pintos booting with 3,968 kB RAM...
367 pages available in kernel pool.
367 pages available in user pool.
Calibrating timer... 104,755,200 loops/s.
Boot complete.
WHUOS> whoami
20250227
WHUOS> ls
invalid command
WHUOS> exit
Bye!
```



Thank you!

Any questions?

